

Exercise 3

Machine Learning in Graphics & Vision

Driton Goxhufi 4233242 driton.goxhufi@student.uni-tuebingen.de
Damir Ravlija 5503184 damir.ravlija@student.uni-tuebingen.de

1 Task 1


- (a) Our implementation does achieve 97.05% accuracy after 15 epochs.

training accuracy

training loss

validation accuracy

validation accuracy



According to http://rodrigob.github.io/are_we_there_yet/build/classification_datasets_results.html#4d4e495354 the paper "Regularization of Neural Networks using DropConnect" in 2013 achieved a result of 0.21% errorrate, hence an accuracy value of 99.79%.

- (b) from torchsummary import summary
summary(model, (1,28,28))

Layer (type)	Output Shape	Param #
--------------	--------------	---------

=====

Conv2d-1 [-1, 32, 28, 28] 320

ReLU-2 [-1, 32, 28, 28] 0

Conv2d-3 [-1, 32, 28, 28] 9,248

ReLU-4 [-1, 32, 28, 28] 0

Conv2d-5 [-1, 32, 28, 28] 9,248

Linear-6 [-1, 256] 6,422,784

ReLU-7 [-1, 256] 0

Linear-8 [-1, 10] 2,570

=====

Total params: 6,444,170

Trainable params: 6,444,170

Non-trainable params: 0

Input size (MB): 0.00
Forward/backward pass size (MB): 0.96
Params size (MB): 24.58
Estimated Total Size (MB): 25.55

Output of convolutional layers can be calculated with this formula:

$$\lfloor (W - K + 2P) / S \rfloor + 1$$

where W is the width(or height if the input is squared), K is the kernelsize, P is the amount of padding and S the stride.

- (c) After 15 epochs training on the default network with FashionMNIST dataset, the validation(test) accuracy is about 89%.

Console output:

epoch done: 14 accuracy train 0.9082844034830729 accuracy test 0.891600112915039

After tweaking (hyper)parameters of the network:

- (d)