

Exercise 3

Machine Learning in Graphics & Vision

Driton Goxhufi 4233242 driton.goxhufi@student.uni-tuebingen.de
 Damir Ravlija 5503184 damir.ravlija@student.uni-tuebingen.de

Task 1


- (a) Our implementation does achieve 97.05% accuracy after 15 epochs.

training accuracy

training loss

validation accuracy

validation accuracy



According to http://rodrigob.github.io/are_we_there_yet/build/classification_datasets_results.html#4d4e495354 the paper "Regularization of Neural Networks using DropConnect" in 2013 achieved a result of 0.21% errorrate, hence an accuracy value of 99.79%.

- (b) from torchsummary import summary
 summary(model, (1,28,28))

Layer (type) Output Shape Param #

=====

Conv2d-1 [-1, 32, 28, 28] 320

ReLU-2 [-1, 32, 28, 28] 0

Conv2d-3 [-1, 32, 28, 28] 9,248

ReLU-4 [-1, 32, 28, 28] 0

Conv2d-5 [-1, 32, 28, 28] 9,248

Linear-6 [-1, 256] 6,422,784

ReLU-7 [-1, 256] 0

Linear-8 [-1, 10] 2,570

=====

Total params: 6,444,170

Trainable params: 6,444,170

Non-trainable params: 0

Input size (MB): 0.00
 Forward/backward pass size (MB): 0.96
 Params size (MB): 24.58
 Estimated Total Size (MB): 25.55

Output of convolutional layers can be calculated with this formula:

$$\lceil (W - K + 2P) / S \rceil + 1$$

where W is the width(or height if the input is squared), K is the kernelsize, P is the amount of padding and S the stride.

- (c) After 15 epochs training on the default network with FashionMNIST dataset, the validation(test) accuracy is about 89%.

Console output:

epoch done: 14 accuracy train 0.9082844034830729 accuracy test 0.891600112915039

We then increased the size of the network in a fashion similar to VGG network. We used stride 2 to speed up training and decrease the number of features while simultaneously increasing the number of filters. To increase the training speed we added batch normalization layers. The following network had the highest accuracy of 0.9108000755310058 on the validation set (the corresponding accuracy curves are in Figure 1):

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 28, 28]	320
BatchNorm2d-2	[-1, 32, 28, 28]	64
ReLU-3	[-1, 32, 28, 28]	0
Conv2d-4	[-1, 32, 14, 14]	9,248
BatchNorm2d-5	[-1, 32, 14, 14]	64
ReLU-6	[-1, 32, 14, 14]	0
Conv2d-7	[-1, 32, 14, 14]	9,248
BatchNorm2d-8	[-1, 32, 14, 14]	64
ReLU-9	[-1, 32, 14, 14]	0
Conv2d-10	[-1, 64, 7, 7]	18,496
BatchNorm2d-11	[-1, 64, 7, 7]	128
ReLU-12	[-1, 64, 7, 7]	0
Conv2d-13	[-1, 64, 7, 7]	36,928
BatchNorm2d-14	[-1, 64, 7, 7]	128
ReLU-15	[-1, 64, 7, 7]	0
Conv2d-16	[-1, 64, 7, 7]	36,928
BatchNorm2d-17	[-1, 64, 7, 7]	128
ReLU-18	[-1, 64, 7, 7]	0
Conv2d-19	[-1, 128, 4, 4]	73,856
BatchNorm2d-20	[-1, 128, 4, 4]	256

ReLU-21	$[-1, 128, 4, 4]$	0	
Conv2d-22	$[-1, 128, 4, 4]$	147,584	
BatchNorm2d-23	$[-1, 128, 4, 4]$		256
ReLU-24	$[-1, 128, 4, 4]$	0	
Conv2d-25	$[-1, 128, 4, 4]$	147,584	
BatchNorm2d-26	$[-1, 128, 4, 4]$		256
ReLU-27	$[-1, 128, 4, 4]$	0	
Conv2d-28	$[-1, 256, 2, 2]$	295,168	
BatchNorm2d-29	$[-1, 256, 2, 2]$		512
ReLU-30	$[-1, 256, 2, 2]$	0	
Conv2d-31	$[-1, 512, 2, 2]$	1,180,160	
BatchNorm2d-32	$[-1, 512, 2, 2]$		1,024
ReLU-33	$[-1, 512, 2, 2]$	0	
Conv2d-34	$[-1, 512, 2, 2]$	2,359,808	
BatchNorm2d-35	$[-1, 512, 2, 2]$		1,024
ReLU-36	$[-1, 512, 2, 2]$	0	
Conv2d-37	$[-1, 512, 2, 2]$	2,359,808	
BatchNorm2d-38	$[-1, 512, 2, 2]$		1,024
ReLU-39	$[-1, 512, 2, 2]$	0	
Conv2d-40	$[-1, 1024, 1, 1]$	4,719,616	
BatchNorm2d-41	$[-1, 1024, 1, 1]$		2,048
ReLU-42	$[-1, 1024, 1, 1]$	0	
Conv2d-43	$[-1, 1024, 1, 1]$	9,438,208	
BatchNorm2d-44	$[-1, 1024, 1, 1]$		2,048
ReLU-45	$[-1, 1024, 1, 1]$	0	
Conv2d-46	$[-1, 1024, 1, 1]$	9,438,208	
BatchNorm2d-47	$[-1, 1024, 1, 1]$		2,048
ReLU-48	$[-1, 1024, 1, 1]$	0	
Linear-49	$[-1, 10]$	10,250	

Total params: 30,292,490
 Trainable params: 30,292,490
 Non-trainable params: 0


Input size (MB): 0.00
 Forward/backward pass size (MB): 1.45
 Params size (MB): 115.56
 Estimated Total Size (MB): 117.01

As we can see from the plots (Figure 1) it seems that our network overfits the data and it could benefit from early stopping.


(d)

Task 2

- (a) Visualization of the image pairs. Every noisy image is located in front of the original image (Figure (2)).



(a) Training accuracy curve



(b) Validation accuracy curve

Figure 1: Plot of results from task 3.1.c)




Figure 2: Sample of image pairs contained in the NoisyFashionMNIST