



# MACHINE LEARNING IN GRAPHICS & VISION

## EXERCISE 3

Release date: Thursday, 21 May 2020 - **Deadline for Homework: Wed, 10 June 2020 - 21:00**

### Excercises

Please do not use jupyter-notebooks for these exercises. Create a new file (e.g., `task_3_1_a.py`) for each (sub-)task and provide your obtained text-output as a comment in the particular file when necessary. Submit your overall assignment in the form of a PDF report along with your code files (Preferably in a .zip folder).

#### 3.1 Hands-On Neural Network Layers (3+2+2+3 Points)

Commonly used layers in deep neural networks are: 2D-convolutions (`torch.nn.Conv2d`) and ReLU (`torch.nn.ReLU`). The file `task_3_1_a.py` contains a simple CNN architecture for classification applied to MNIST digits.

- a) There are missing parts and two bugs in the provided implementation template. Complete the implementation and remove the bugs. Plot the accuracy curve against the epochs. Which validation-accuracy does your implementation achieves after 15 epochs after completing the implementation? (**Hint:** Use a TensorBoard-screenshot (preferred) or you can print your data in a log file.) What is the state-of-the-art accuracy on classifying MNIST digits reported in the literature?
- b) Provide a model summary which includes the list of the names and output shapes for each used `torch.nn` layers in your Model. How to compute the output shape of (`torch.nn.Conv2d`) with zero padding (valid) for an input of shape [B, T, T, C] and stride 2 kernel-size 3. (**Hint:** You can use the `torchsummary` package.)
- c) Run the network from a) on the Fashion MNIST dataset. Which validation-accuracy do you get after 15 training epochs? Try to improve the network by creating more layers or tweaking the hyper-parameters and explain your findings.
- d) We provide 10 images in the template from the Fashion MNIST dataset. Your task is to create an offline predictor function (`task_3_1_d.py`) which runs inference on these images. Which labels are predicted by your network for these images? (**Hint:** You may want to look at how to save and load a model for inference on PyTorch [tutorials](#).)

### 3.2 Data Augmentation and Denoising (4 + 4 + 2 Points)

Augmenting inputs helps when the amount of training data is relatively low. Changing the data flow can help to train a neural network on a different task without a time-consuming data collection.

- a) Create a dataflow class which “yields” FashionMNIST image pairs (noisy, clean) using `torch.utils.data.Dataset` as presented in the dataflow tutorial. Therefore, fill-out the missing parts in `task_3_2_a.py:NoisyFashionMNIST()` and add uniform distribution noise  $(0,1)$  to the input image. Debug your dataflow by visualizing these image pairs and save the visualization as a JPEG image.
- b) A classic task is denoising images. We will use a simple encoder-decoder network  $\pi_\theta$  with trainable parameters  $\theta$  to solve this task. The network  $\pi_\theta$  is trained on noisy/clean image pairs to predicts a single denoised image.

$$\text{inference: } \pi_\theta: \underbrace{\mathbb{R}^{H \times W \times C}}_{\text{noisy image}} \rightarrow \underbrace{\mathbb{R}^{H \times W \times C}}_{\text{predicted denoised image}}, \quad X \mapsto \pi_\theta(X), \quad (1)$$

$$\text{training: } \min_{\theta} \|\pi_\theta(X) - X_{\text{gt}}\|_2. \quad (2)$$

Hereby, for  $\pi_\theta$  we propose the following structure:

- 2D-convolution layer (16 filters, kernel size 3x3, stride 2 and padding=1)
- 2D-convolution layer (16 filters, kernel size 3x3, stride 2 and padding=1)
- 2D-convolution layer (16 filters, kernel size 3x3, stride 1 and padding=1)
- transposed 2D-convolution layer (16 filters, kernel size 3x3, stride 2 and padding=1)
- transposed 2D-convolution layer (16 filters, kernel size 3x3, stride 2 and padding=1)
- 2D-convolution layer (16 filters, kernel size 3x3, stride 1 and padding=1 no activation!)

You are free to implement a more sophisticated version. Please do not forget to add a non-linear activation function (we suggest using `torch.nn.ReLU`). We ask you to implement this network structure in the provided template and train the model. Which *Mean Absolute Error (MAE)*-validation loss do you observe. The Tensorboard will show some results on the training example. Interpret, these results.

- c) A common trick is to add a skip-connection by training a the network  $\pi'_\theta$  which has the following structure:

$$\pi'_\theta: \mathbb{R}^{H \times W \times C} \rightarrow \mathbb{R}^{H \times W \times C}, \quad X \mapsto X + \pi_\theta(X). \quad (3)$$

Train the modified network  $\pi'_\theta$  separably and compare the result again  $\pi_\theta$  in terms of *MAE*-validation loss. Explain possible benefits from adding such a skip-connection.

### Hints

- For Tensorboard implementation, you can refer to the official [Pytorch tutorial](#) on Tensorboard basics.
- Look in the tutorials for PyTorch layers explanation and reference links (Conv2D, Maxpool2D, etc.).