

MACHINE LEARNING IN GRAPHICS & VISION

EXERCISE 5

Release date: Thu, 24. June 2020 - **Deadline for Homework: Wed, 08. July 2020 - 21:00**

So far, we have mainly considered *supervised learning* where we provide both inputs and desired outputs to the model during training. In this exercise, we consider *unsupervised learning* where the goal is to make a machine learning model learn about patterns in the data on its own. We consider both a classical approach, *Principal Component Analysis* and a more modern approach, *Variational Autoencoders*¹ (VAEs), in this exercise. Both models can be considered as probabilistic models that are trained by maximizing (a lower bound to) the log-likelihood

$$\frac{1}{N} \sum_{i=1}^N \log p_{\theta}(x_i) \quad (1)$$

of a parametric probabilistic model $p_{\theta}(x)$. For PCA, $p_{\theta}(x)$ is a multivariate Gaussian distribution. For VAEs, $p_{\theta}(x)$ is given by the marginal distribution of a latent variable model $p_{\theta}(x, z) = p_{\theta}(x | z)p(z)$.

Exercises

For this exercise you need to submit a **.zip** file containing your report as a **.pdf** file and your code (namely the files `report.pdf`, `utils.py`, `exercise_5_1.py`, `exercise_5_2.py`). Make sure that your report is self contained and includes all your results and visualizations. **You will only get points for results that are included in the report. Please use the provided code templates for these exercises. Do not use jupyter-notebooks.**

For the visualizations, you can use the functions provided in the `utils.py` script.

5.1 Principal Component Analysis (2 + 2 + 2 + 2 + 2 Points)

In this exercise, we apply PCA to the Fashion-MNIST dataset.

- a) Compute the mean vector of the Fashion-MNIST dataset and apply PCA to the *centered* dataset (subtract the mean from the samples before applying PCA). Fill in the function

```
compute_pca(x)
```

where x is $N \times K$ input array for N data points of dimensionality K . Test your implementation and visualize the first 5 principal components.

- b) Plot the percentage of explained variance over the number of principal components. How many components do you need to explain 50%, 90%, 95%, 99% of the variance?
- c) We can use PCA to compress test data. For each sample in the *test set* compute the coefficients belonging to each principal component. Reconstruct each test sample using only the first 5 principal components from the previous exercise (don't forget to center your data using the mean vector). Report the averaged mean-squared error (MSE) over all test samples. Visualize the reconstructions of the first five test samples. What is the compression ratio?

¹Kingma, Diederik P., and Max Welling. "Auto-encoding variational bayes." Proceedings of the 2nd International Conference on Learning Representations (ICLR), 2014.

- d) We can use PCA to sample from the probabilistic model obtained by fitting a Gaussian distribution to the training set. To do so, we sample $\epsilon \in \mathbb{R}^K$ from $\mathcal{N}(0, I)$ and then compute $x = \mu + \sum_{k=1}^K \sqrt{s_k} v_k \epsilon_k$ where μ denotes the mean vector, s_k denote the k^{th} singular values of the covariance matrix and v_k the k^{th} principal component. Draw 5 samples from this probabilistic model and visualize them. Do they look realistic?
- e) Fashion-MNIST contains 10 different classes and is therefore highly varied. Can you obtain better results by applying PCA only to the *Sneaker*-subset (with class id 7) of this dataset? Return your code on this subset and report your results.

5.2 Variational Autoencoder (4 + 4 + 2 Points)

We now want to improve our model by applying a Variational Autoencoder to the dataset. Recall that VAEs are trained by maximizing the so-called *Evidence Lower Bound* (ELBO) to the log-likelihood which is given by

$$\frac{1}{N} \sum_{i=1}^N \log p_{\theta}(x_i) \geq \frac{1}{N} \sum_{i=1}^N \mathbb{E}_{q_{\phi}(z|x_i)} [\log p_{\theta}(x_i | z)] - \text{KL}(q_{\phi}(z | x_i) | p(z)) \quad (2)$$

with an additional inference model (the encoder) $q_{\phi}(z | x)$. Here, the (negative of the) first term can be interpreted as a reconstruction loss and the second one as a regularizer. Usually, $p(z)$ is a standard Gaussian distribution and $q_{\phi}(z | x)$ is a Gaussian distribution with diagonal covariance matrix. Recall that the KL-divergence between a Gaussian distribution $q_{\phi}(z | x)$ with mean μ and diagonal covariance matrix with diagonal entries σ^2 is given by

$$\text{KL}(q(z | x) | p(z)) = \frac{1}{2} \sum_k (-\log(\sigma_k^2) - 1 + \sigma_k^2 + \mu_k^2) \quad (3)$$

- a) Complete the function

```
compute_kl
```

which computes the KL-divergence between a multivariate Gaussian distribution with diagonal covariance matrix and a unit Gaussian distribution. Test your implementation by calling

```
test_kl
```

with different values for μ and σ .

- b) Implement a fully connected VAE where both the encoder and decoder have two hidden layers with 512 units each. Use RELU-activations and a latent dimension of 5. Use cross-entropy (with logits) as reconstruction error² and train the model for 10 epochs. How do the different losses behave during training? What do you observe?
- c) Use the trained VAE to compress the test data using the inference model $q_{\phi}(z | x)$. Compute the average mean-squared error for the corresponding reconstructions (using $p_{\theta}(x | z)$) and visualize the first 5 reconstructions. How do the reconstructions compare to the ones obtained using PCA? Sample from the learned distribution and visualize 5 samples. Do they look realistic?

²This corresponds to using independent Bernoulli variables for the conditional distribution $p_{\theta}(x | z)$